

# CKAD\_beta\_exam

## Basic Information:

- total 120 minutes
- total 20 programmatic questions
- questions may have different weights and different difficulty
- exam results will be delivered in three business days, otherwise you have to contact to linux foundation by yourself

## Questions:

### **.bashrc**

```
source <(kubectl completion bash)

alias kl="kubectl "
alias klc="kubectl apply -f "
alias klpod="kubectl get pods --all-namespaces -o wide"
alias kldep="kubectl get deployment --all-namespaces -o wide"
alias klall="kubectl get all --all-namespaces -o wide"
```

### **create namespace**

```
# delete namespace
kubectl delete namespace redis;
kubectl delete namespace pod-resources;
kubectl delete namespace kdpd00201;
kubectl delete namespace kdpd00202;
kubectl delete namespace kdpd00203;

# create namespace
kubectl create namespace redis;
kubectl create namespace pod-resources;
kubectl create namespace kdpd00201;
kubectl create namespace kdpd00202;
kubectl delete namespace kdpd00203;
```

1	<p>Set configuration context: \$ kubectl config use-context k8s</p> <p>A web application requires a specific version of Redis to be used as a cache.</p> <p>Create a Pod with the following characteristics, and leave it running when complete:</p> <p>The Pod should run in the <b>redis</b> namespace. The namespace has already been created</p> <p>The name of the Pod should be <b>redis-server</b></p> <p>Use the <b>lfccncf/redis</b> image with the <b>3.2</b> tag</p> <p>Expose <b>containerPort 6379</b></p> <p>Question weight: 13%</p>	<p><a href="https://kubernetes.io/docs/user-guide/walkthrough/k8s201/">https://kubernetes.io/docs/user-guide/walkthrough/k8s201/</a></p> <p># delete labels, update others</p> <pre>\$ kubectl create namespace reids # for test  # p1-pod.yaml apiVersion: v1 kind: Pod metadata:   name: redis-server   namespace: redis spec:   containers:   - name: redis     image: lfccncf/redis:3.2   ports:   - containerPort: 80  # check \$ kubectl --namespace=redis get po redis-server NAME          READY    STATUS    RESTARTS   AGE redis-server  1/1      Running   0           1m  \$ kubectl --namespace=redis log redis-server</pre>
---	---	---

2	<p>Set configuration context: \$ kubectl config use-context k8s</p> <p>You are tasked to create a Secret and consume the Secret in a Pod using environment variables as follows:</p> <p>Create a Secret named <b>app-secret</b> with a key/value pair: <b>key2/value1</b></p> <p>Start an nginx Pod named <b>nginx-secret</b> using container image <b>nginx</b>,</p> <p>and add an environment variable exposing the value of the secret key <b>key2</b>,</p> <p>using <b>BEST_VARIABLE</b> as the name for the environment variable inside the Pod</p> <p>Question weight: 5%</p>	<p><a href="https://kubernetes.io/docs/tasks/inject-data-application/distribute-credentials-secure/">https://kubernetes.io/docs/tasks/inject-data-application/distribute-credentials-secure/</a></p> <pre>\$ kubectl create secret generic app-secret --from-literal=key2='value1'  # p2-pod.yaml apiVersion: v1 kind: Pod metadata:   name: nginx-secret spec:   containers:   - name: envvars-test-container     image: nginx     env:     - name: BEST_VARIABLE       valueFrom:         secretKeyRef:           name: app-secret           key: key2  # check \$ kubectl get po nginx-secret NAME          READY   STATUS    RESTARTS   AGE nginx-secret  1/1     Running   0           1m \$ kubectl exec -it nginx-secret -- /bin/bash \$ printenv   grep BEST_VARIABLE BEST_VARIABLE=value1</pre>
---	---	--

3	<p>Set configuration context: \$ kubectl config use-context k8s</p> <p>You are required to create a Pod that requests a certain amount of CPU and memory,</p> <p>so it gets scheduled to a node that has those resources available.</p> <p>Create a Pod named <b>nginx-resources</b> in the <b>pod-resources</b> Namespace</p> <p>that <b>requests</b> a minimum of <b>200m</b> CPU and <b>2Gi</b> memory for its container</p> <p>The Pod should use the <b>nginx</b> image</p> <p>The pod-resources Namespace has already been created</p> <p>Question weight: 4%</p>	<p><a href="https://kubernetes.io/docs/tasks/configure-pod-container/assign-memory-resource/#specify-a-memory-request-and-a-m">https://kubernetes.io/docs/tasks/configure-pod-container/assign-memory-resource/#specify-a-memory-request-and-a-m</a></p> <pre>\$ kubectl create namespace pod-resources # for test  # p3-pod.yaml apiVersion: v1 kind: Pod metadata:   name: nginx-resources   namespace: pod-resources spec:   containers:   - name: memory-demo-ctr     image: nginx     resources:       requests:         cpu: "200m"         memory: "2Gi"  # check \$ kubectl --namespace=pod-resources get po NAME                READY    STATUS    RESTARTS   AGE nginx-resources     1/1      Running   0           1m  \$ kubectl --namespace=pod-resources describe po nginx-resources   grep "cpu\ memory" memory-demo-ctr:   cpu:          200m   memory:       2Gi</pre>
---	---	---

<div>4</div> <div>Set configuration context: \$ kubectl config use-context k8s</div> <div>You are tasked to create a ConfigMap and consume the ConfigMap in a Pod using a volume mount.</div> <div>Please complete the following:</div> <div>Create a ConfigMap named <b>some-config</b> containing the key/value pair: <b>key3/value4</b> Start a Pod named <b>nginx-configmap</b> containing a single container using the <b>nginx</b> image,</div> <div>and mount the key you just created into the Pod under directory <b>/some/other/path</b></div> <div>Question weight: 5%</div>	<div><a href="https://kubernetes.io/docs/tasks/configure-pod-container/configure-pod-configmap/#add-configmap-data-to-a-specific-pod">https://kubernetes.io/docs/tasks/configure-pod-container/configure-pod-configmap/#add-configmap-data-to-a-specific-pod</a></div> <div><pre>\$ echo "key3=value4" &gt; env.txt \$ kubectl create configmap some-config --from-env-file=env.txt  # p4-pod.yaml apiVersion: v1 kind: Pod metadata:   name: nginx-configmap spec:   containers:     - name: test-container       image: nginx       command: [ "/bin/sh", "-c", "ls /some/other/path/"       volumeMounts:         - name: config-volume           mountPath: /some/other/path   volumes:     - name: config-volume       configMap:         name: some-config   restartPolicy: Never  # check \$ kubectl get po nginx-configmap NAME                READY    STATUS    RESTARTS   AGE nginx-configmap     0/1      Completed  0           57s \$ kubectl logs nginx-configmap key3</pre></div>
---	---

5	<p>Set configuration context: \$ kubectl config use-context k8s</p> <p>A Pod is running on the cluster but it is not responding.</p> <p>The desired behavior is to have Kubernetes restart the pod when an endpoint returns an HTTP 500 on the /healthz endpoint.</p> <p>The service, liveness-pod, should never send traffic to the Pod while it is failing. Please complete the following:</p> <p>The application has an endpoint, /started, that will indicate if it can accept traffic by returning an HTTP 200.</p> <p>If the endpoint returns an HTTP 500, the application has not yet finished initialization</p> <p>The application has another endpoint /healthz that will indicate if the application is still working as expected by returning an HTTP 200.</p> <p>If the endpoint returns an HTTP 500 the application is no longer responsive</p> <p>Configure the liveness-pod Pod provided to use these endpoints</p> <p>The probes should use port 8080</p> <p>Question weight: 2%</p>	
6	<p>Set configuration context: \$ kubectl config use-context k8s</p> <p>You sometimes need to observe a Pod's logs, and write those logs to a file for further analysis. Please complete the following:</p> <p>Deploy the <b>counterapp</b> Pod to the cluster using the provided YAML spec file at <b>/opt/KDOB00201/counterapp.yaml</b></p> <p>Retrieve all currently available application logs from the running Pod and store them in the file <b>/opt/KDOB00201/log_output.txt</b>,</p> <p>which has already been created</p> <p>Question weight: 4%</p>	<pre>\$ view /opt/KDOB00201/counterapp.yaml \$ kubectl get po counterapp \$ kubectl logs -p counterapp &gt; /opt/KDOB00201/log_output.txt \$ less /opt/KDOB00201/log_output.txt</pre>

7	<p>set configuration context: \$ kubectl config use-context k8s</p> <p>Create a new <b>Deployment</b> for running nginx with the following parameters:</p> <p>Run the Deployment in the <b>kdpd00201</b> Namespace. The Namespace has already been created</p> <p>Name the Deployment <b>frontend</b> and configure with <b>3</b> replicas</p> <p>Configure the Pod with a container image of <b>lfccncf/nginx:1.12.2</b></p> <p>Set an environment variable of <b>NGINX_PORT=8001</b> and also <b>expose that port</b> for the container above</p> <p>Question weight: 5%</p>	<p><a href="https://kubernetes.io/docs/user-guide/walkthrough/k8s201/#deployment-management">https://kubernetes.io/docs/user-guide/walkthrough/k8s201/#deployment-management</a></p> <pre>\$ kubectl create namespace kdpd00201  # vim p6-deploy.yaml apiVersion: apps/v1 kind: Deployment metadata:   name: frontend   namespace: kdpd00201 spec:   selector:     matchLabels:       app: nginx   replicas: 3   template:     metadata:       labels:         app: nginx     spec:       containers:         - name: nginx           image: lfccncf/nginx:1.12.2           env:             - name: NGINX_PORT               value: "8001"  # check \$ kubectl --namespace=kdpd00201 get all \$ kubectl --namespace=kdpd00201 describe pod/frontend-9bc77c74d-fl267</pre>
8	<p>Set configuration context: \$ kubectl config use-context k8s</p> <p>As a Kubernetes application developer you will often find yourself needing to update a running application. Please complete the following:</p> <p>Update the <b>web1</b> <b>Deployment</b> in the <b>kdpd00202</b> Namespace with a <b>maxSurge</b> of 4 and a <b>maxUnavailable</b> of 2</p> <p>Perform a rolling update of the <b>web1</b> Deployment, changing the <b>lfccncf/nginx</b> image version to <b>1.13</b></p> <p>Roll back the <b>web1</b> Deployment to the previous version</p> <p>Question weight: 6%</p>	

## test file

```
$ kubectl create namespace kdpd00202

# test file
$ vim web1-deploy.yaml

apiVersion: apps/v1
kind: Deployment
metadata:
  name: web1
  namespace: kdpd00202
spec:
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%
  selector:
    matchLabels:
      app: nginx
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: lfcncf/nginx:1.12.2

# check
$ kubectl --namespace=kdpd00202 get deploy web1
```



		<pre> \$ kubectl --namespace=kdpd00202 edit deployment web1 # change maxSurge of 4 and a maxUnavailable of 2 deployment "web1" edited  \$ kubectl --namespace=kdpd00202 set image deployment/web1 nginx=lfccncf/nginx:1.13 deployment "web1" image updated  \$ kubectl --namespace=kdpd00202 get deploy web1 -o wide NAME           DESIRED   CURRENT   UP-TO-DATE   AVAILABLE CONTAINERS     IMAGES                SELECTOR web1           3         3         3            3 nginx          lfccncf/nginx:1.13   app=nginx  \$ kubectl --namespace=kdpd00202 rollout undo deployment/web1 deployment "web1" rolled back  \$ kubectl --namespace=kdpd00202 get deploy web1 -o wide NAME           DESIRED   CURRENT   UP-TO-DATE   AVAILABLE CONTAINERS     IMAGES                SELECTOR web1           3         3         3            3 nginx          lfccncf/nginx:1.12.2 app=nginx </pre>
9	find maximum cpu usage pod name in a namespace	use top?
10	config service account, and use that account to create a pod	
11	create deployment and service	
12	logs	
13		

14	<p>Given a container that writes a log file in format A and a container that converts log files from format A to format B,</p> <p>create a Deployment that runs both containers such that the log files from the first container are converted by the second container,</p> <p>emitting logs in format B.</p> <p>Create a Deployment named <b>deployment-007</b> in the default namespace, that:</p> <ul style="list-style-type: none"><li>Includes a primary <b>lfcncf/busybox:1</b> container, named <b>logger-zen</b></li><li>Includes a sidecar <b>lfcncf/fluentd:v0.12</b> container, named <b>adaptor-123</b></li></ul> <p>Mounts a shared volume <b>/tmp/log</b> on both containers, which does not persist when the pod is deleted</p> <p>Instructs the logger-zen container to run the command <b>while true; do echo i luv cncf &gt;&gt; /tmp/log/input.log; sleep 10; done,</b></p> <p>which should output logs to <b>/tmp/log/input.log</b> in plain text format, with example values:</p> <pre>i luv cncf i luv cncf i luv cncf</pre> <p>The adaptor-123 sidecar container should read <b>/tmp/log/input.log</b> and output the data to <b>/tmp/log/output.*</b> in Fluentd JSON format.</p> <p>Note that no knowledge of Fluentd is required to complete this task: all you will need</p> <p>Question weight: 8%</p>	<p><a href="https://kubernetes.io/docs/concepts/cluster-administration/logging/#sidecar-container-with-a-logging-agent">https://kubernetes.io/docs/concepts/cluster-administration/logging/#sidecar-container-with-a-logging-agent</a></p> <div></div>
----	--	--

<div>15</div> <div>Set configuration context: \$ kubectl config use-context k8s</div> <div>Developers occasionally need to submit Pods that run periodically. Follow the steps below to create a Pod</div> <div>that will start at a predetermined time and which runs to completion only once each time it is started:</div> <div>Create a YAML formatted Kubernetes manifest <code>/opt/KDPD00301/periodic.yaml</code> that runs the following shell command:</div> <div><code>date</code> in a single <code>busybox</code> container</div> <div>The command should run every <code>minute</code> and must complete within 10 seconds or be terminated by Kubernetes.</div> <div>The <code>CronJob</code> name and <code>container</code> name should both be <code>hello</code></div> <div>Create the resource in the above manifest and verify that the job executes successfully at least once</div> <div>Question weight: 4%</div>	<div><a href="https://kubernetes.io/docs/concepts/workloads/controllers/cron-jobs/">https://kubernetes.io/docs/concepts/workloads/controllers/cron-jobs/</a></div> <div><pre># cronjob.yaml apiVersion: batch/v1beta1 kind: CronJob metadata:   name: hello spec:   activeDeadlineSeconds: 10 # add   schedule: "*/1 * * * *"   jobTemplate:     spec:       template:         spec:           containers:             - name: hello               image: busybox               args:                 - /bin/sh                 - -c                 - date;           restartPolicy: OnFailure  # check</pre></div> <div>kubectl explain cronjob --api-version=batch/v1beta1</div>
--	---

16	<p>Set configuration context: \$ kubectl config use-context k8s</p> <p>You have been tasked with scaling an existing Deployment for availability,</p> <p>and creating a Service to expose the Deployment within your infrastructure.</p> <p>Start with the Deployment named <b>kdsn00101-deployment</b> which has already been deployed to the namespace <b>kdsn00101</b>. Edit it to:</p> <p>Add the <b>role=webFrontEnd</b> <b>key/value</b> label to the Pod template metadata to identify the pod for the Service definition Have <b>4</b> replicas Next, create and deploy in namespace <b>kdsn00101</b> a Service that accomplishes the following:</p> <p>Exposes the Service on TCP port <b>8080</b> Is mapped to the pods defined by the specification of <b>kdsn00101-deployment</b> Is of type <b>NodePort</b> Has a name of <b>cherry</b> Question weight: 6%</p>	<p><a href="https://kubernetes.io/docs/concepts/services-networking/service/">https://kubernetes.io/docs/concepts/services-networking/service/</a></p> <p><a href="https://kubernetes.io/docs/concepts/services-networking/connect-applications-service/">https://kubernetes.io/docs/concepts/services-networking/connect-applications-service/</a></p> <pre> kubectl --namespace=kdsn00101 apply -f dep.yaml  kubectl --namespace=kdsn00101 edit deployment kdsn00101-deployment role:webFrontEnd # two loc replicas 4  kind: Service apiVersion: v1 metadata:   name:  cherry # add spec:   type: NodePort   selector:     role: webFrontEnd # add   ports:     - protocol: TCP       nodePort: 30000       port: 8080       targetPort: 80 </pre> <p>kubectl --namespace=kdsn00101 apply -f svc.yaml</p>
17	<p>Set configuration context: \$ kubectl config use-context nk8s</p> <p>You have rolled out a new Pod to your infrastructure and now you need to allow it to communicate with the <b>www</b> and <b>api</b> pods but nothing else.</p> <p>Given the running pod <b>kdsn00201-newpod</b> edit it to use a network policy that will allow it to send and receive traffic only to and from the <b>www</b> and <b>api</b> Pods.</p> <p>All work on this item should be conducted in the <b>kdsn00201</b> namespace.</p> <p>All required <b>NetworkPolicy</b> resources are already created and ready for use as appropriate.</p> <p>You should not create, modify or delete any network policies whilst completing this item.</p> <p>Question weight: 6%</p>	<p><a href="https://kubernetes.io/docs/concepts/services-networking/network-policies/">https://kubernetes.io/docs/concepts/services-networking/network-policies/</a></p> <p><a href="https://kubernetes.io/docs/tasks/administer-cluster/declare-network-policy/">https://kubernetes.io/docs/tasks/administer-cluster/declare-network-policy/</a></p> <p>change pods labels</p>

18	<p>Set configuration context: \$ kubectl config use-context dk8s</p> <p>A user has reported an application is unreachable due to a failing <b>livenessProbe</b>.</p> <p>The associated deployment could be running in any of the following namespaces:</p> <p><b>qa</b> <b>test</b> <b>production</b> <b>alan</b></p> <p>Perform the following tasks:</p> <p>Find the broken Pod and store its <b>name</b> and <b>namespace</b> to <b>/opt/KDOB00401/broken.txt</b> in the format <b>&lt;namespace&gt;/&lt;pod&gt;</b>.</p> <p>The output file has already been created Store the associated <b>error</b> events to a file <b>/opt/KDOB00401/error.txt</b>. The output file has already been created Fix the issue Question weight: 6%</p>	<a href="https://kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-probes/">https://kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-probes/</a>
19	<p>Set configuration context: \$ kubectl config use-context sk8s</p> <p>A project that you are working on has a requirement for persistent data to be available. To facilitate this, perform the following tasks:</p> <p>Create a file on <b>node sk8s-node-0</b> at <b>/opt/KDSP00101/data/index.html</b> with the content <b>PW=MyStuff!</b> Create a PersistentVolume <b>name d task-pv-volume</b> using <b>hostPath</b> and allocate <b>2Gi</b> to it,</p> <p>specifying that the volume is at <b>/opt/KDSP00101/data</b> on the cluster's Node.</p> <p>The configuration should specify the access mode of <b>ReadWriteOnce</b>. It should define the <b>StorageClass</b> name storage for the PersistentVolume,</p> <p>which will be used to bind PersistentVolumeClaim requests to this PersistentVolume Create a PersistentVolumeClaim named <b>task-pv-claim</b> that requests a volume of at least <b>100Mi</b> and specifies an access mode of <b>ReadWriteOnce</b> Create a Pod that uses the PersistentVolumeClaim as a volume with a label <b>app: my-storage-app</b> mounting the resulting volume to</p> <p>a mountPath <b>/usr/share/nginx/html</b> inside the Pod Hint: You can access sk8s-node-0 by issuing the following command:</p> <p>ssh sk8s-node-0</p>	<a href="https://kubernetes.io/docs/tasks/configure-pod-container/configure-persistent-volume-storage/">https://kubernetes.io/docs/tasks/configure-pod-container/configure-persistent-volume-storage/</a>

Ensure that you return to the base node (with hostname node-1) once you have completed your work on sk8s-node-0

Question weight: 6%

```
kind: PersistentVolume
apiVersion: v1
metadata:
  name: task-pv-volume
  labels:
    type: local
spec:
  storageClassName: k8s #update
  capacity:
    storage: 2Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/home/liguohui9527/q19/data" #update

kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: task-pv-claim
spec:
  storageClassName: k8s #update
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 200Mi #update

kind: Pod
apiVersion: v1
metadata:
  name: task-pv-pod
  labels:
    #update
  app: my-storage-app #update
spec:
  volumes:
    - name: task-pv-storage
      persistentVolumeClaim:
        claimName: task-pv-claim
  containers:
    - name: task-pv-container
      image: nginx
      ports:
        - containerPort: 80
          name: "http-server"
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: task-pv-storage
```

```
kubectl delete pod task-pv-pod;
kubectl delete pvc task-pv-claim;
kubectl delete pv task-pv-volume;
```

```
kubectl apply -f pv.yaml;
kubectl get pv task-pv-volume;
kubectl apply -f pvc.yaml;
kubectl get pv task-pv-volume;
kubectl get pvc task-pv-claim;
kubectl apply -f pod.yaml;
kubectl get pod task-pv-pod;
kubectl describe po task-pv-pod;
```

Lessons:

- DO NOT USE WORK COMPUTER! Or at least turn off firewall!
- very practical questions and quite like real-world problem solving
- Time is short and stressful. If you cannot finish a question within 10 minutes, then just skip to the next one
- create a new folder for each question, and double check your solution with the question before you execute your command
- you can view the Kubernetes document during the exam, better to memorize the key operations instead of reading the documents in the exam
- at least 2 ~ 3 months dedicated working hours on Kubernetes will help you to have a good understanding on all the concepts and pass the exam